



# Kokkos, a sustainable performance portability layer

co-developed by CExA, a project for Computing at Exascale on Accelerators at CEA



Adrien Taberner  
Developer



Dr. Ansar Calloo  
Group leader, Senior Developer



Dr. Cédric Chevallier

Group leader, Senior developer



Dr. Hariprasad Kannan  
Developer



Dr. Mathieu Lobet  
Developer



Dr. Paul Zehner  
Developer



Dr. Rémi Baron  
Senior Developer



Dr. Thomas Padioleau  
Senior developer, architect



Dr. Yuuichi Asahi  
Senior developer

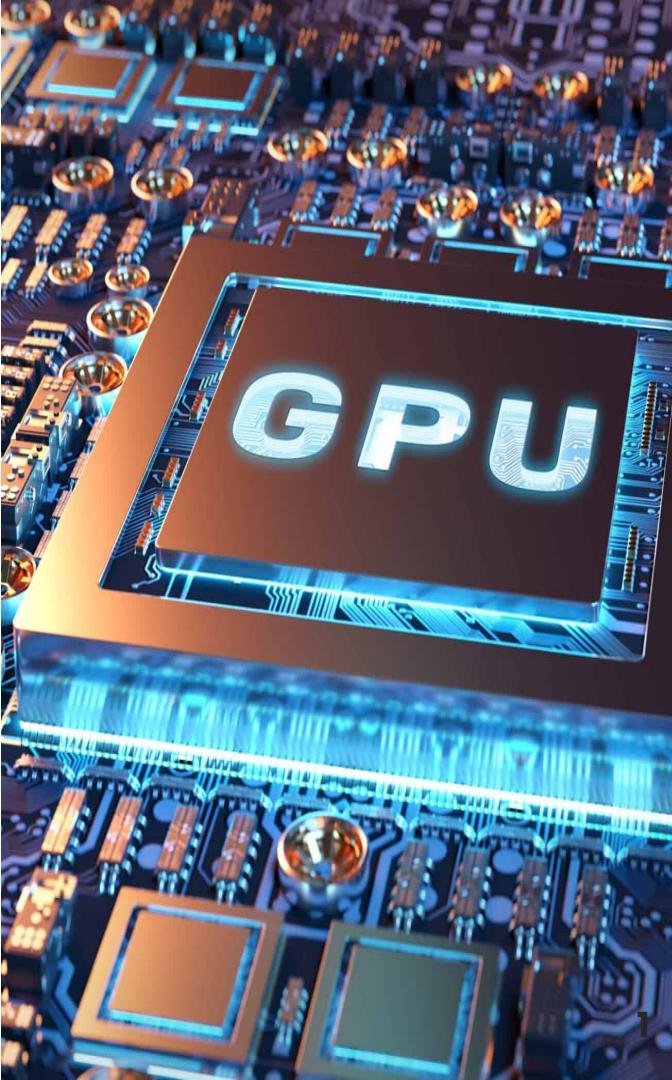


Paul Gannay  
Developer



Trévis Morvany  
Developer

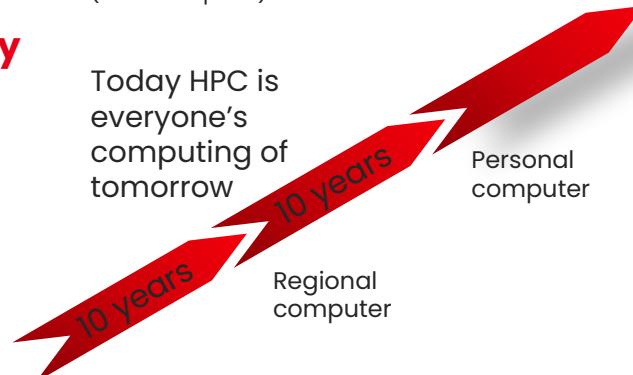
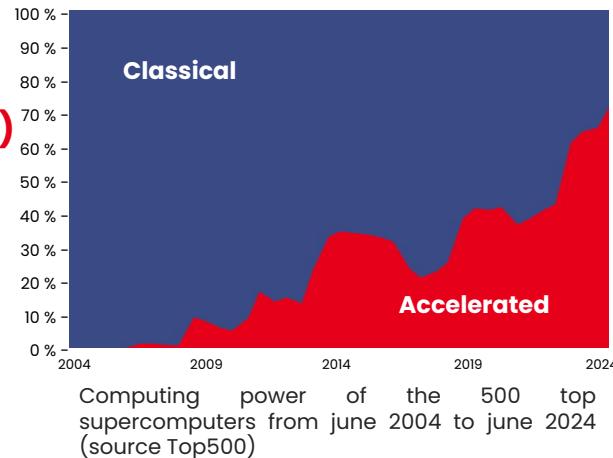
ORAP Forum  
May 19<sup>th</sup> 2025  
Julien Bigot, the CExA, Kokkos & HPSF teams





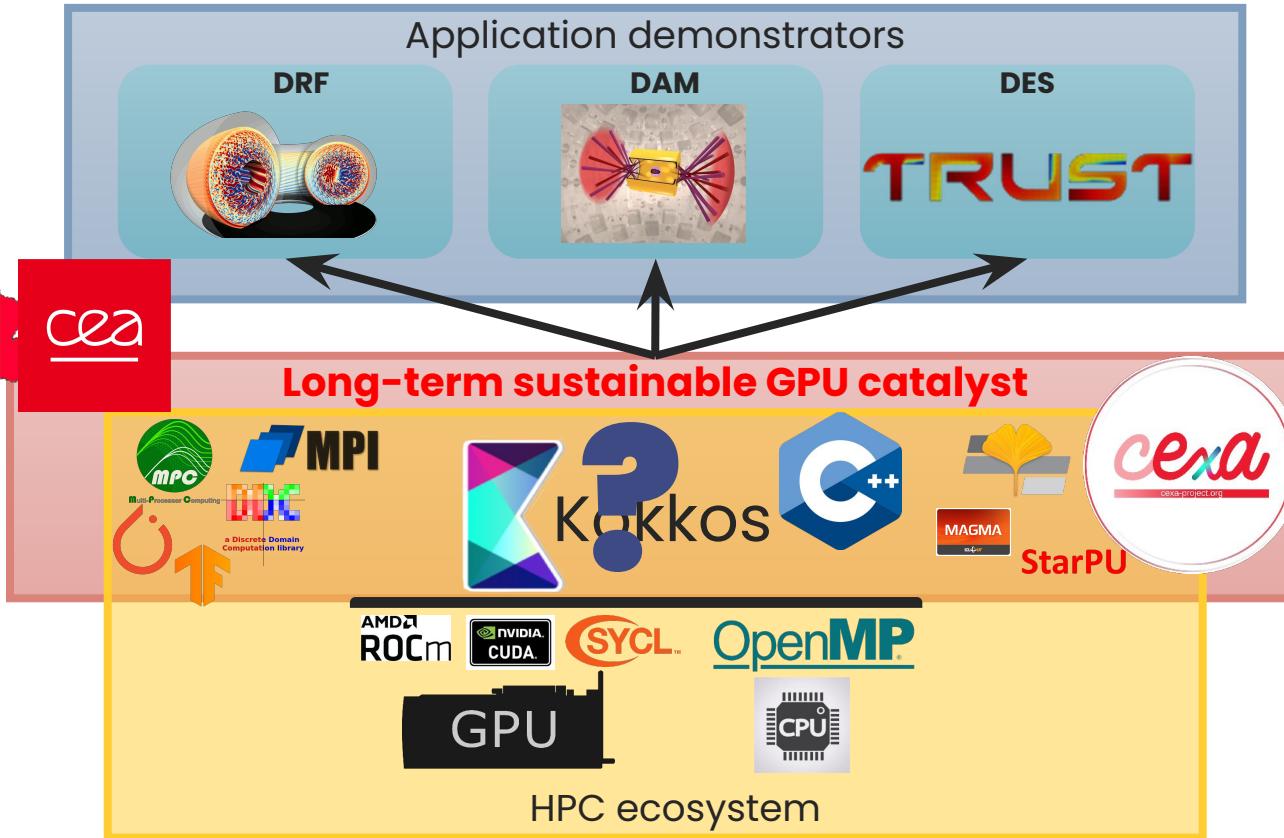
# Context 1+ year ago

- HPC is a tool in a wide range of domains at CEA
- We just entered the **Exascale** era, that means **GPU (>82% Top500)**
  - European pre-Exascale systems: Mix of **AMD & Nvidia**
  - US Exascale systems: **AMD & Intel**
  - First European Exascale machines are coming
    - *Jupiter* at Jülich (Germany) => **Nvidia & Rhea**
    - *Alice Recoque* at **CEA/TGCC** (**call still open**)
  - Need to re-develop applications with **Performance portability**
- GPU programming models: **software catalysts**
  - France and Europe: great research but no production tool
- A **need** for a long-term sustainable solution
  - **Adapted** to our hardware and software specificities
  - Where we can have **trust** in the roadmap





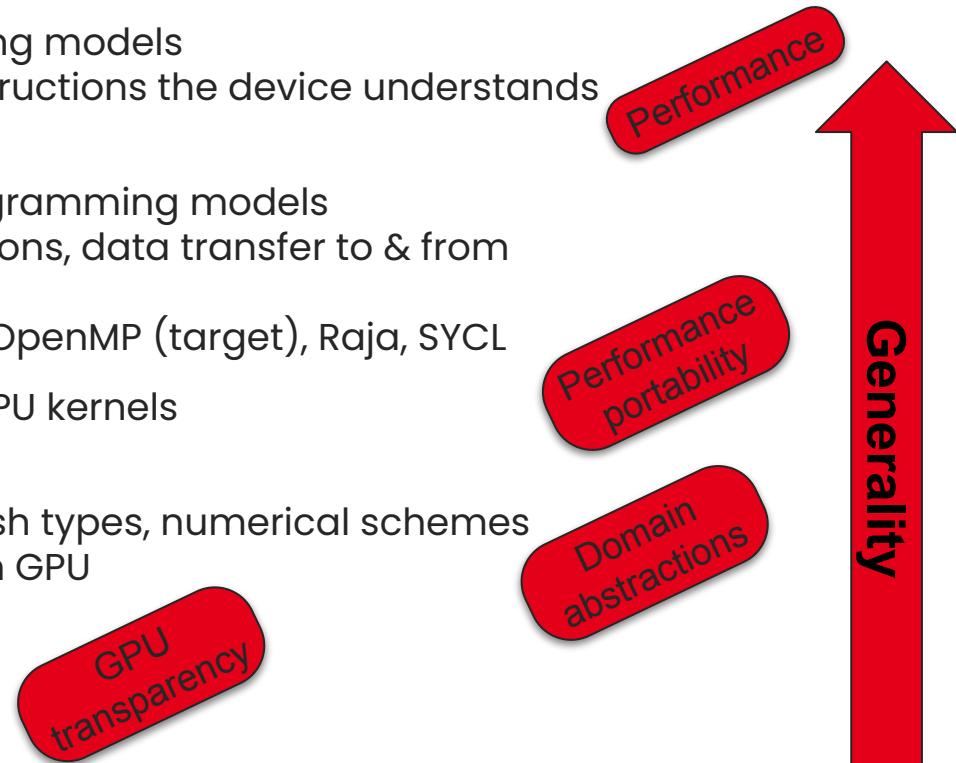
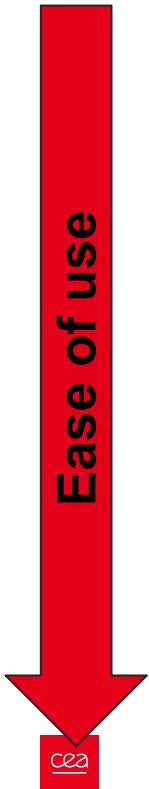
# CExA project: goals





# How to write code for GPU?

- Low-level, assembly-style programming models
  - Nearly manipulate the actual instructions the device understands
  - E.g. HSA, Level Zero, PTX, Spir-V , ...
- General-purpose, imperative GPU programming models
  - Manipulate parallel loops, reductions, data transfer to & from device
  - E.g. Cuda, HIP, Kokkos, OpenACC, OpenMP (target), Raja, SYCL
- Combination & assembly of existing GPU kernels
  - Pytorch, StarPU, etc...
- Application framework for specific mesh types, numerical schemes
  - Use domain-specific concepts on GPU
- Pre-written GPU libraries
  - just call them from CPU
  - Neural Networks, Linear Algebra, ...





# How to write code for GPU?

- Low-level, assembly-style programming models
  - Nearly manipulate the actual instructions the device understands
  - E.g. HSA, Level Zero, PTX, Spir-V , ...
- General-purpose, imperative GPU programming models
  - Manipulate parallel loops, reductions, data transfer to & from device
  - E.g. Cuda, HIP, Kokkos, OpenACC, OpenMP (target), Raja, SYCL
- Combination & assembly of existing GPU kernels
  - Pytorch, StarPU, etc...
- Application framework for specific mesh types, numerical schemes
  - Use domain-specific concepts on GPU
- Pre-written GPU libraries
  - just call them from CPU
  - Neural Networks, Linear Algebra, ...

Ease of use

Generality

GPU transparency

Performance portability

Domain abstractions

Performance



# How to write code for GPU?

- Low-level, assembly-style programming models
  - Nearly manipulate the actual instructions the device understands
  - E.g. HSA, Level Zero, PTX, Spir-V , ...
- General-purpose, imperative GPU programming models
  - Manipulate parallel loops, reductions, data transfer to & from device
  - E.g. **Cuda, HIP, Kokkos, OpenACC, OpenMP (target), Raja, SYCL**
- Combination & assembly of existing GPU kernels
  - Pytorch, StarPU, etc...
- Application framework for specific mesh types, numerical schemes
  - Use domain-specific concepts on GPU
- Pre-written GPU libraries
  - just call them from CPU
  - Neural Networks, Linear Algebra, ...

Ease of use

cea

Performance

Performance portability

Domain abstractions

GPU transparency

Generality



# Available solutions

- Cuda
- HIP
- Kokkos
- OpenACC
- OpenMP (target)
- Raja
- SYCL
  - OneAPI/DPC++
  - AdaptiveC++/OpenSYCL/hipSYCL



# Available solutions

- Cuda
- HIP
- Kokkos
- OpenACC
- OpenMP (target)
- **Raja**
- SYCL
  - OneAPI/DPC++
  - **AdaptiveC++/OpenSYCL/hipSYCL**
- **Production grade, with public support**



# Available solutions

- Cuda
- HIP
- Kokkos
- OpenACC
- OpenMP (target)
- Raja
- SYCL
  - OneAPI/DPC++
  - AdaptiveC++/OpenSYCL/hipSYCL
- Production grade, with public support
- **Vendor neutral**



# Available solutions

- Cuda
- HIP
- **Kokkos**
- OpenACC
- **OpenMP (target)**
- Raja
- SYCL
  - OneAPI/DPC++
  - AdaptiveC++/OpenSYCL/hipSYCL
- Production grade, with public support
- Vendor neutral



# OpenMP & Kokkos : the simplest GPU loop

```
for (int j = 0 ; j < Nj ; ++j) {  
    // [...]  
}
```

Sequential



```
#pragma omp teams distribute parallel for  
for (int j = 0 ; j < Nj ; ++j) {  
    // [...]  
}
```

OpenMP Target

```
parallel_for(Nj, KOKKOS_LAMBDA(int j) {  
    // [...]  
});
```

Kokkos

Execute in **parallel**, on a separate GPU thread each,

the same workload **[ . . . ]**

identified by a unique identifier **j**

**Nj** times between 0 and Nj-1



# OpenMP & Kokkos : memory transfer

```
double* x = malloc(Ni*sizeof(double));
double* y = malloc(Nj*sizeof(double));
double* A = omp_target_alloc(
    Ni*Nj*sizeof(double),
    omp_get_initial_device());

#pragma omp target data \
    map(to: x[0:Ni]) \
    map(from: y[0:Nj])
{
#pragma omp teams distribute parallel for
for (int j = 0 ; j < Nj ; ++j) {
    for (int i = 0 ; i < Ni ; ++i) {
        y[j] += x[i] * A[j*Ni+i];
    }
}
```

OpenMP Target

Copy x to GPU from device before kernel  
and y from GPU to device after kernel

Keep A on the device

```
View<double*, Kokkos::HostSpace> x(Ni);
View<double*, Kokkos::HostSpace> y(Nj);
View<double*> A(Nj, Ni);

{
auto dx = create_mirror_view_and_copy(dev, x);
auto dy = create_mirror_view(dev, y);
parallel_for(Nj, KOKKOS_LAMBDA(int j) {
    for (int i = 0 ; i < Ni ; ++i) {
        dy(j) += dx(i) * A(j,i);
    }
});
deep_copy(y, dy);
}
```

Kokkos



# Compilation

## OpenMP Target

- Use an OpenMP compiler
  - Compatible with the target construct
  - Compatible with the hardware you target
- Each vendor provides its own OpenMP compiler
  - Usually based on LLVM infra
- Default Clang/LLVM & GCC also try to support this
  - For some hardware
  - With variable performance

## Kokkos

- A C++ template library
  - No direct code generation, rely on vendors C++-like languages
- Multiple “backends”, selection at compile time
  - Each vendor’s preferred toolchain
  - OpenMP, Cuda, OneAPI, HIP, ...
- Maximum 3 backends enabled at once
  - Serial backend
  - 1 Host parallel backend (openmp)
  - 1 Device parallel backend (cuda, HIP, Sycl)



# Available solutions

- Cuda
- HIP
- **Kokkos**
- OpenACC
- **OpenMP (target)**
- Raja
- SYCL
  - OneAPI/DPC++
  - AdaptiveC++/OpenSYCL/hipSYCL
- Production grade, with public support
- Vendor neutral



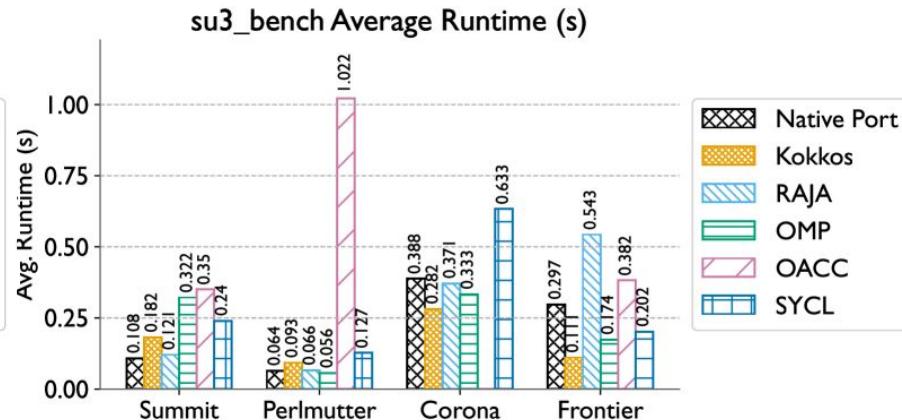
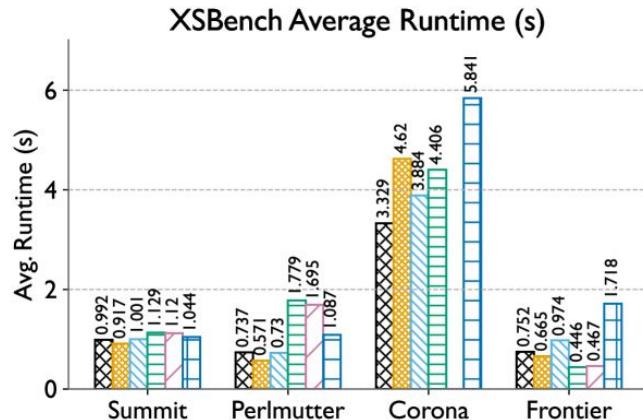
# Available solutions

- Cuda
- HIP
- **Kokkos**
- OpenACC
- **OpenMP (target)**
- Raja
- SYCL
  - OneAPI/DPC++
  - AdaptiveC++/OpenSYCL/hipSYCL

- Production grade, with public support
- Vendor neutral
- **Annotations**
  - Works best with **imperative languages**: C, Fortran, ...
  - **Compiler integration**: potential for additional optimizations
  - **Seq. first**, requires to re-design applications for GPU
- **Library**
  - Suited to language with deep **encapsulation**: C++, ...
  - On top of vendor toolchains: easier to port to **new hardware**
  - **GPU first**, requires to re-write applications for GPU



# And what about performance?



Average runtime of the XSBench (left) and su3\_bench (right) proxy apps across all platforms and programming models.

**Lower is better.**

## An Evaluative Comparison of Performance Portability across GPU Programming Models

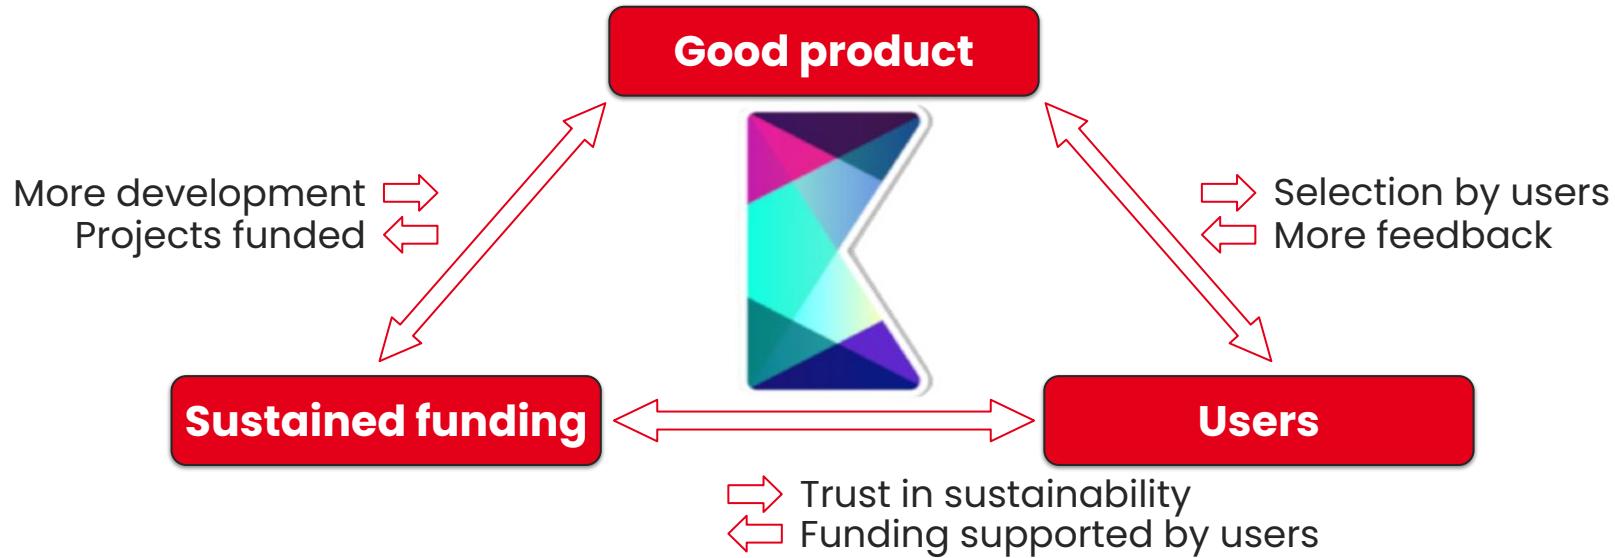
Joshua H. Davis<sup>2</sup>, Pranav Sivaraman<sup>2</sup>, Isaac Minn<sup>2</sup>, Konstantinos Parasyris<sup>1</sup>, Harshitha Menon<sup>1</sup>, Giorgis Georgakoudis<sup>1</sup>, Abhinav Bhatele<sup>2</sup>

<sup>2</sup>Department of Computer Science, University of Maryland

<sup>1</sup>Center for Applied Scientific Computing, Lawrence Livermore National Laboratory



# Kokkos at the center of a virtuous cycle

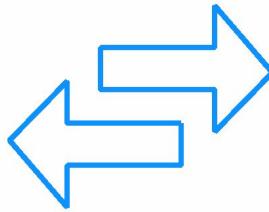


**There is strength in numbers:  
collaboration on core products is good for everyone**

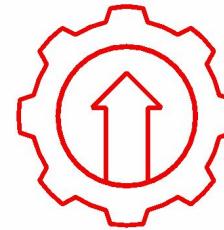
# Kokkos, a HPSF project



Performance

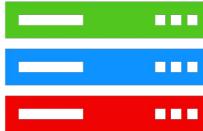


Portability



Productivity

1. A neutral hub for open source, high performance software
2. Supports projects that advance portable software for diverse hardware by:
  - o Increasing adoption
  - o Aiding community growth
  - o Enabling development efforts
3. Lower barriers to productive use of today's and future high performance computing systems



# HPSF

HIGH PERFORMANCE  
SOFTWARE FOUNDATION

Fund & vote

## Members



Premier

Hewlett Packard  
Enterprise

Sandia  
National  
Laboratories



Lawrence  
Livermore  
National  
Laboratory

aws  
Hewlett Packard  
Enterprise

Sandia  
National  
Laboratories

Lawrence Livermore  
National Laboratory

cea

kokkos

Governing board

Participate & vote



WG&s & committees

Technical Advisory Council

Outreach

Diversity

CI &  
Testing

Events

Tools

...



Projects

Spack

Charliecloud

AMReX

HPCToolkit

APPTAINER

EAS

TRILINOS

kokkos

OpenCHAMI

VISKORES

HPX

cea

WarpX

ENVIRONMENT  
MODULES

CHAPEL

General

AMD

Argonne  
NATIONAL LABORATORY

arm

cea

intel

kitware  
Delivering Innovation

NVIDIA

Los Alamos  
NATIONAL LABORATORY  
EST. 1943

OAK RIDGE  
National Laboratory

Riken  
R-CCS

Associate

BERKELEY LAB

सी.डैक  
CDAC

JÜLICH  
Forschungszentrum  
Tennessee  
TECH

IIC

MARYLAND

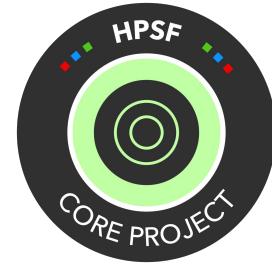
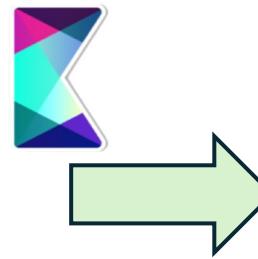
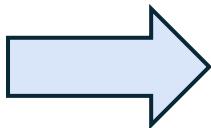
UNIVERSITY OF  
OREGON

UTokyo

der Bundeswehr  
Universität München



# The TAC has established a project lifecycle as a path to sustainability



## Emerging

- Committed to open governance
- Working towards best practices
- Important projects for the HPC ecosystem

## Established

- Wide usage by at least 3 orgs of sufficient size and scope
- Steady commits from at least one organization
- Robust development practices

## Core

- Used commonly in production environments
- Steady commits from *more than one* organization
- Large, well-established project communities
- Sustainable cycle of development and maintenance



**HPSF**

HIGH PERFORMANCE  
SOFTWARE FOUNDATION

# Neutrality through HPSF



1. Sustaining OSS projects requires a community
2. Building a community requires trust
  - o Projects will continue to be available
  - o Projects are usable by anyone
  - o No one organization can control the direction of the project
  - o Projects are open to new contributors and new ideas
3. Trust gets us users; some users become contributors
4. Neutral, open governance ensures that we can build the broadest possible communities



# The CExA project

“adopt and adapt” strategy based on  Kokkos

Kokkos : **a strong technical basis**

- A software architecture ready for the future
- Mature, free, libre, and open-source
- An **independent foundation** to own the product
  - HPSF under the Linux Foundation
- A **standardisation** effort in **ISO C++**
  - A **stepping stone** one step ahead toward **parallel C++**



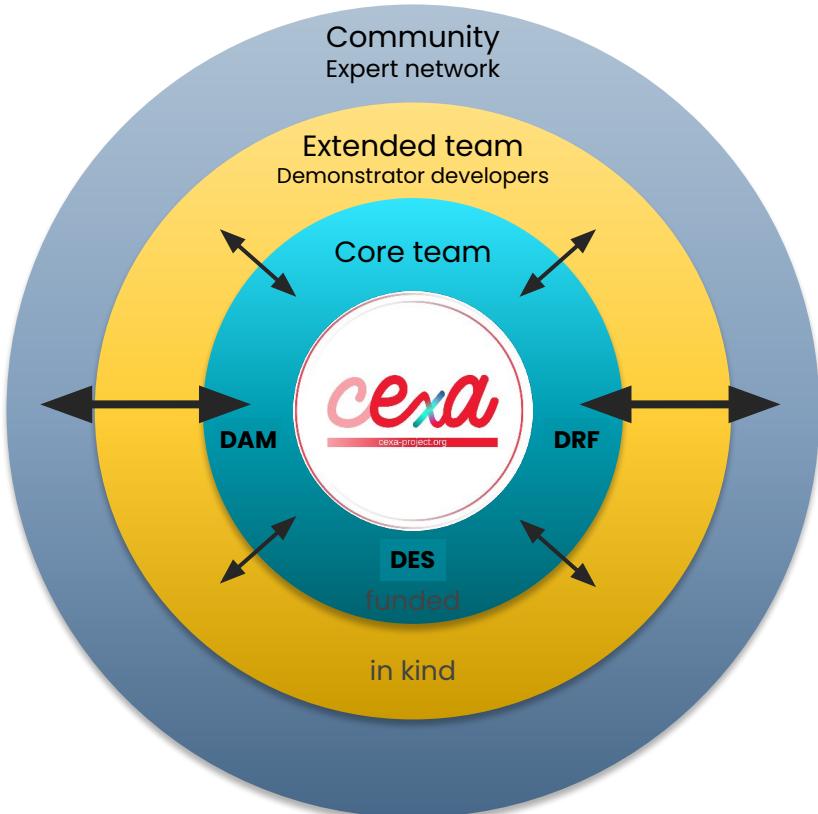
Some **adaptations required**

- For European **hardware**
  - There is no real hardware sovereignty without software sovereignty
- For **applications** from CEA, France and Europe
  - Take our specificities into account





# CExA project in practice



## ■ Core team

- Management, implementation and dissemination
- 12 researchers from all over CEA
- 6 dedicated recruitments
  - **1 as a permanent researcher !**
- Funding for 2 or 3 hires expected every year

## ■ Extended team

- Demonstrator developers
  - Not funded
  - Find their own interest in the participation
  - 2-3 new demonstrators every year

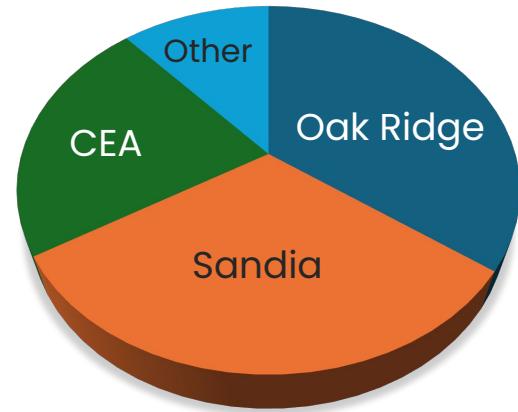
## ■ Community

- Federation of an **expert network**
- Co-design of **CExA**:
  - Identification of needs
  - Usage of **CExA** in applications
- Priority target for **dissemination**
- **Sustainability** of the work



# Kokkos development today

## Primary teams



## Contributions & support



- Number of commits by institution
- In the last 6 month



# Kokkos parallel patterns

```
parallel_for(Nj, KOKKOS_LAMBDA(int j) {  
    // [...]  
}) ;
```



# Kokkos parallel patterns

```
parallel_for(Nj, KOKKOS_LAMBDA(int j) {  
    // [...]  
});
```

```
parallel_reduce(Nj, KOKKOS_LAMBDA(int j, double& accumulator) {  
    // [...]  
    accumulator += /* [...] */ ;  
}, result);
```

```
parallel_scan(Nj, KOKKOS_LAMBDA(int j, double& result, bool isfinal)  
{  
    // [...]  
    accumulator += /* [...] */ ;  
    if(is_final) {  
        // [...]  
    }  
}, result);
```

- For
  - independent iterations
- Reduce
  - Accumulate into a single value
- Scan
  - N independent prefix reduction



# Kokkos parallel patterns: easy debug

```
parallel_for("loop1", Nj, KOKKOS_LAMBDA(int j) {  
    // [...]  
});
```

- Naming loops ease debugging & profiling
- Integrated with kokkos-specific tools
- Get a trace with names includes
- Get a name in debug messages
- Omitted in the presentation, but a good practice overall



# Kokkos parallel patterns: ExecutionSpace

```
parallel_for<ExecutionSpace>(Nj, KOKKOS_LAMBDA(int j) {  
    // [...]  
});
```

- ExecutionSpace defines where to run
  - Cuda, HIP, SYCL, HPX, OpenMP, OpenMPTarget, Threads, Serial
  - 3 exec spaces per execution max: Serial + parallel Host + parallel Device
- Choose where to run at compile time with a #define
  - Usually set from CMake
- 2 predefined aliases are often enough
  - DefaultExecutionSpace: parallel Device, or parallel Host, or Serial
    - Most of the time
  - DefaultHostExecutionSpace: parallel Host, or Serial
    - When using host-only code



# Kokkos parallel patterns: Policies

```
parallel_for(RangePolicy(DefaultExecutionSpace(), 1, Nj, chunk_size), KOKKOS_LAMBDA(int j) {  
    // [...]  
});
```

Beyond simple 1D execution

- Give an instance of ExecutionSpace for multi-GPU or multi-Stream support
- RangePolicy for 1D iteration
  - Begin / end iteration boundaries
  - Chunk\_size hint for improved performance
- MDRange policy for multi-dimensional iterations
  - Multi-D begin / end iteration boundaries
  - Tiling hint hint for improved performance



# Kokkos parallel patterns: hierarchical parallelism

```
parallel_for(TeamPolicy(Nj, team_size), KOKKOS_LAMBDA(const team_handle& team) {  
    // [...]  
    parallel_for(TeamThreadRange(team, Ni, chunk_size), KOKKOS_LAMBDA(int i) {  
        // [...]  
    });  
    // [...]  
});
```

- Default loops can not be nested
- 2-level nesting is supported by teams of threads
  - Matches groups / threads support in GPU
  - But also available on CPU
  - Intermediate (scratch) memory allocation available



# Kokkos parallel patterns are asynchronous

```
parallel_for(Nj, KOKKOS_LAMBDA(int j) {
    // [...]
});
parallel_for(Nj, KOKKOS_LAMBDA(int j) {
    // [...]
});
fence();
```

- Asynchronous execution
- Result visibility is only assured after a fence
- Or between kernels running on the same execution space



# Kokkos views: multi-dimensional arrays

```
View<int**, MemorySpace> my_matrix("matrix", Nx, Ny);
```

- Multi-dimensional arrays
  - Type & dimensionality specified: int\*\* => 2D integer array
  - Dynamic sizes are parameters: Nx, Ny
  - Static sizes are also possible: int\*[4] => 2D array, 4 × dynamic
- Behaves like a C++ shared\_ptr
  - Shared ownership with reference counting (like in python)
- With a name for debugging/profiling
- MemorySpace is part of the type, defaults should be used
  - CudaSpace, CudaHostPinnedSpace, CudaUVMSpace, HIPSpace, HIPHostPinnedSpace, HIPManagedSpace, SYCLDeviceUSMSpace, SYCLHostUSMSpace, SYCLSharedUSMSpace, **HostSpace**, **SharedSpace**, SharedHostPinnedSpace
  - Check of accessibility between MemorySpace & ExecutionSpace



# Kokkos views copies & co.

```
auto dview = subview(oview, pair(start, end), ALL, slice_idx);
```

- Make a new reference to a subset of an existing view
  - Modifying the result modifies the source
  - pair: select a subrange, ALL: keep the dimension, integer: slice the dimension

```
void deep_copy(const ExecSpace &exec_space, const ViewDest &dest, const ViewSrc &src);
```

- Copy data between 2 views
  - Potentially on distinct memory spaces
  - An asynchronous operation

```
auto dview = create_mirror(mspace, a_view); // allocates & copy a new view of same size  
auto dview = create_mirror_view_and_copy(mspace, a_view); // allocates & copy if necessary
```

- Allocates & copy to a new memory space

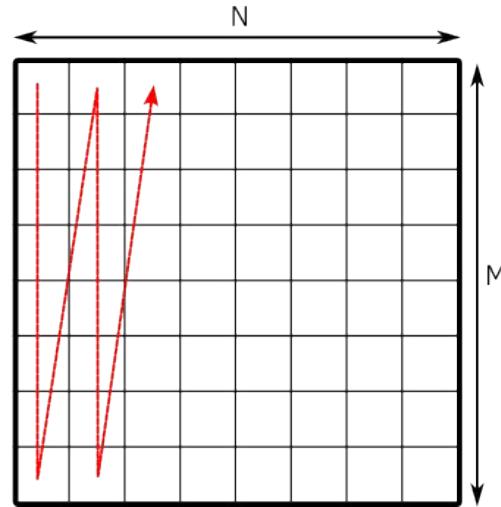


# Kokkos views layout

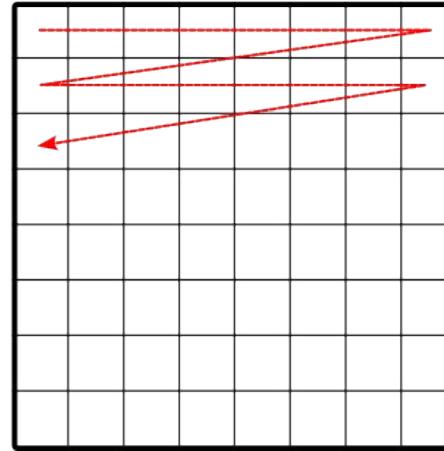
```
View<double**, LayoutLeft> A("A", M, N);
```

- Layout specifies the linearization of multi-D indices into memory
  - LayoutLeft (a.k.a Fortran, default on GPU)
  - LayoutRight (a.k.a C, default on Host)
  - LayoutStride (generic, useful for subviews)

Layout left  
Column major in 2D  
(Fortran)  
Device default layout



Layout right  
Raw major in 2D  
(C, C++, Python, Java)  
Host default layout





# What's in Kokkos (core library)?

Multi-dimensional arrays

- Layout auto change for performance

Parallel patterns w. asynchronous support

- Independent interactions, Reductions, Scans

Iteration strategies

- Tiled, Hierarchical, ...



# What's in Kokkos (core library)?

Multi-dimensional arrays

- Layout auto change for performance

Other containers

- Key-value maps, ScatterView ...

Automatic ref-counted Host/Device memory allocation & management

Host/device memory transfers

Support of "dual" arrays with one version on each side

- Up-to-date tracking & automatic transfers when required

Scratch memory

- Using "team-local" fast memory on the device

Parallel patterns w. asynchronous support

- Independent interactions, Reductions, Scans

Iteration strategies

- Tiled, Hierarchical, ...

Algorithms

- Sorting
- Random number generation
- Many of STL parallel algorithms
- ...

QoL features: portable printf, etc.

Portable atomic operations

SIMD

Coarse & fine-grain tasks

And much more...



# What's new in Kokkos (core library)?

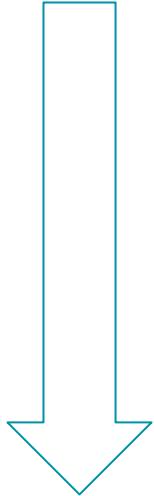
- Multi-GPU support for CUDA, HIP, and SYCL
  - Launch kernels on multiple devices from a single host process
- Many Team-Level Algorithms
  - Extended API with a new overload for team-level support
- Build systems update
  - support linking against build tree & building with hidden visibility
- New Kokkos::View implementation
  - Now mdspan-based
- Improved sorting performance
- Updated SIMD support to align with C++26
- Improved Thread-Safety
  - Kernels can be submitted from multiple threads to the same execution space
- Clarified support for View of Views
  - Introduced a new SequentialHostInit view allocation property
- Improved clang-tidy Conformity



# Kokkos Ecosystem



Kokkos-based applications

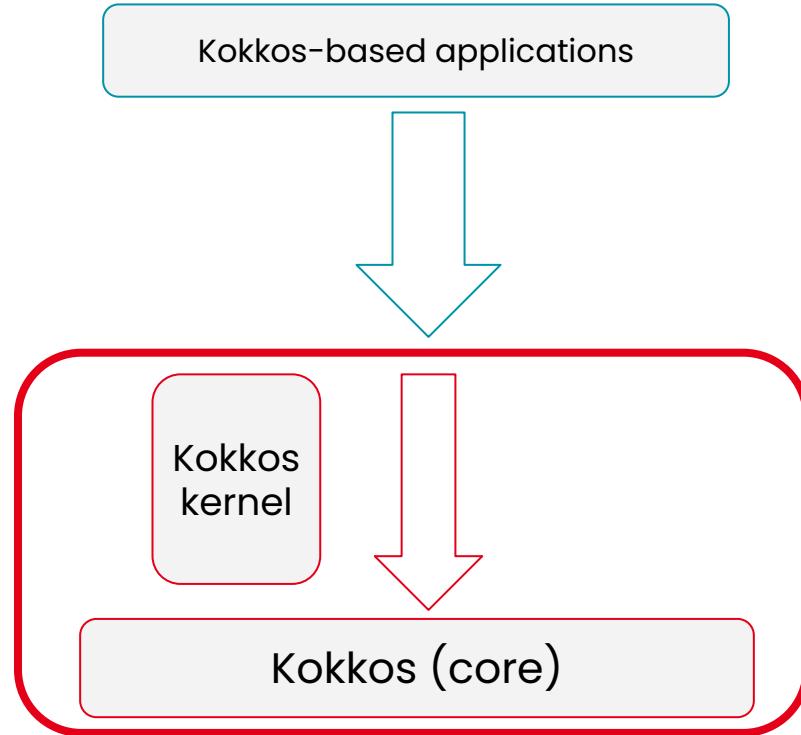


Kokkos (core)



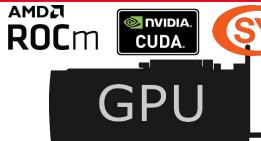
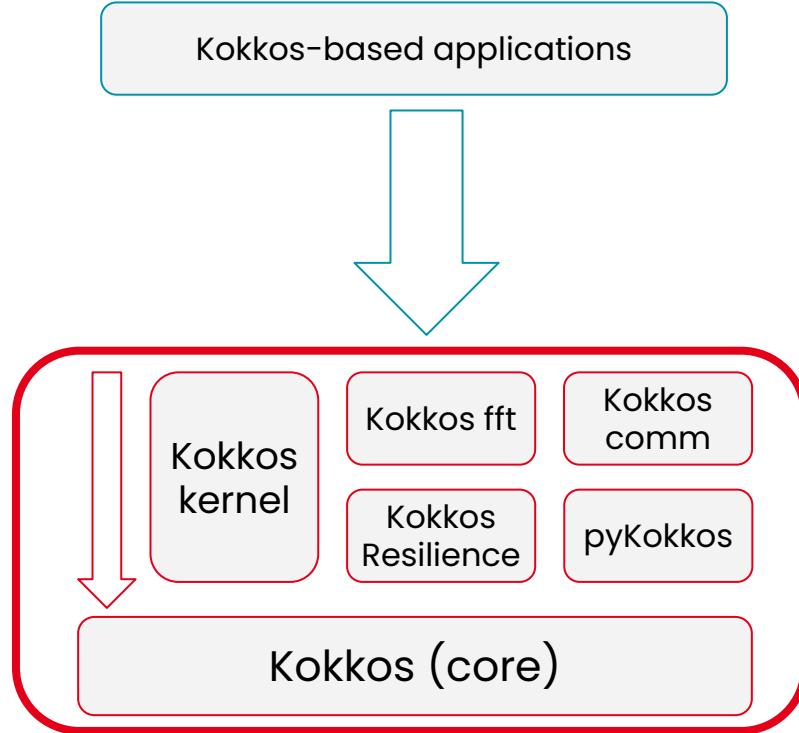


# Kokkos Ecosystem



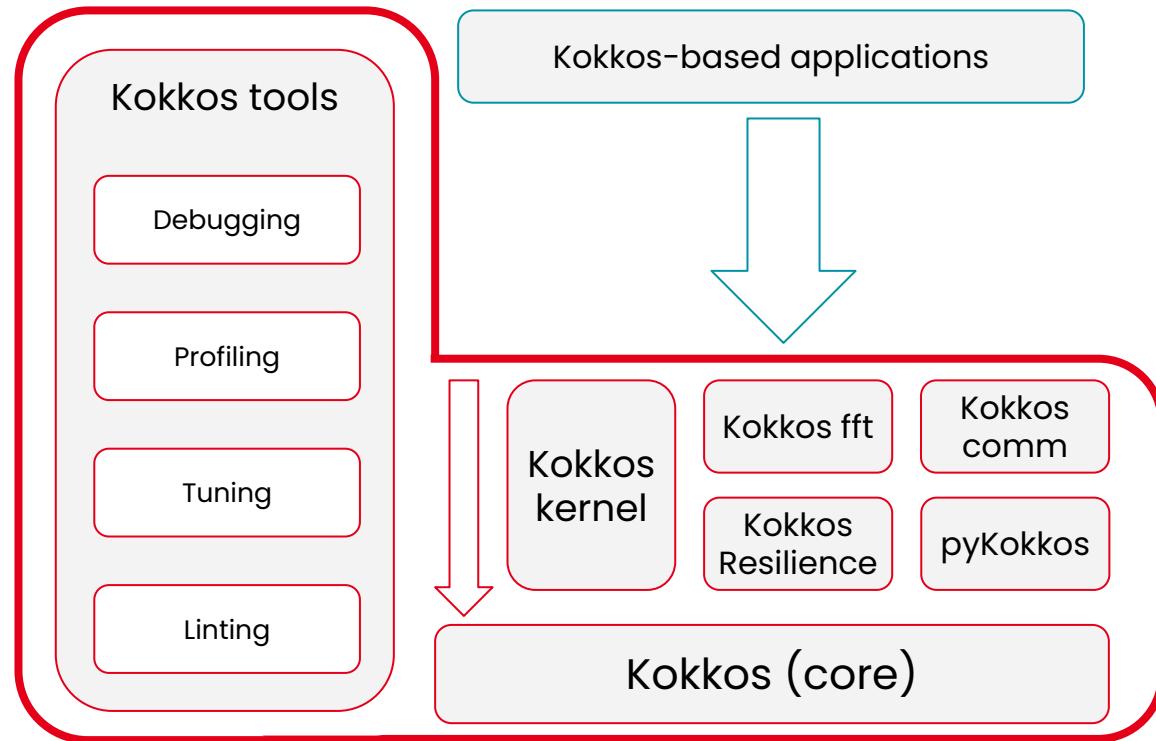


# Kokkos Ecosystem



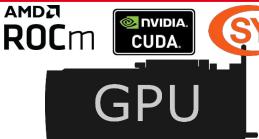
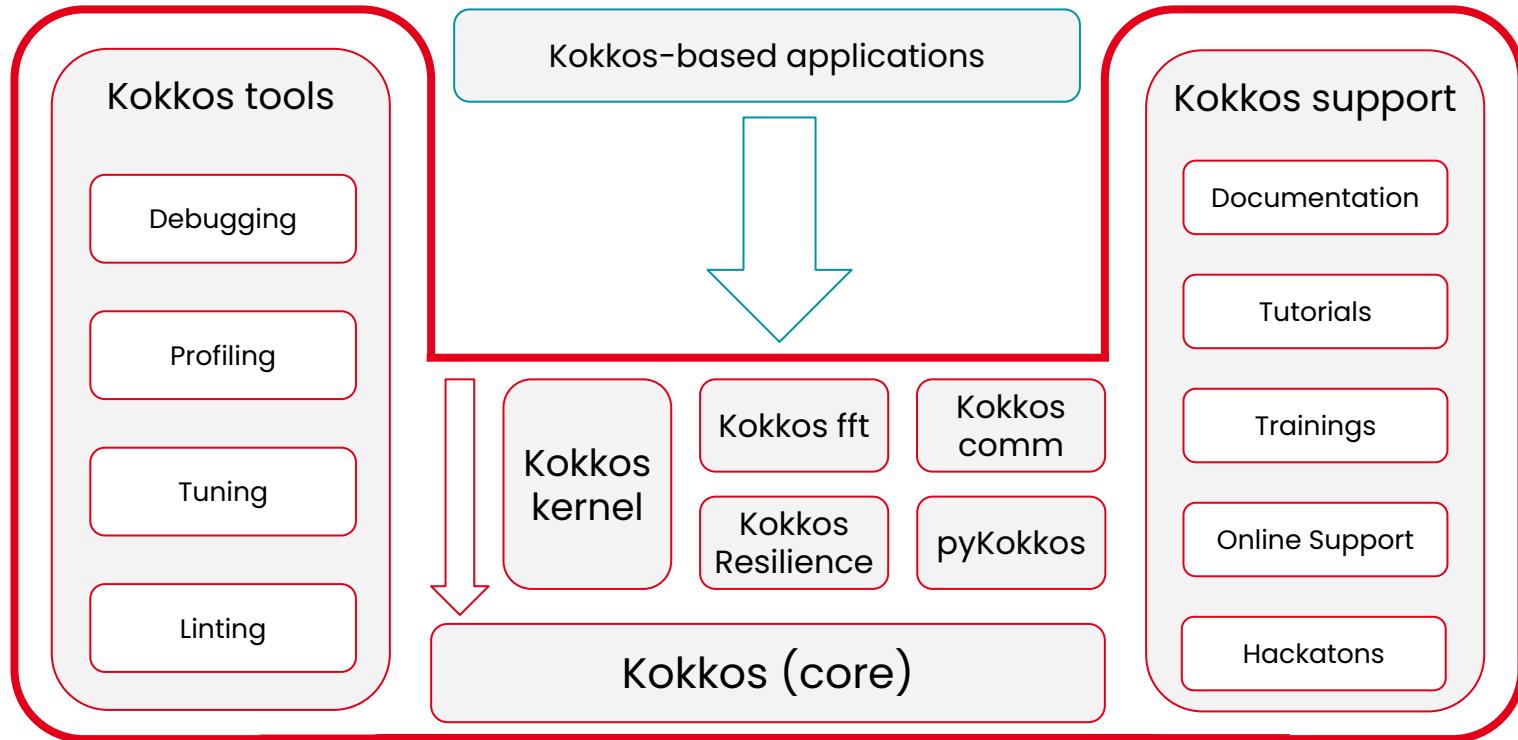


# Kokkos Ecosystem



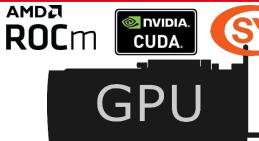
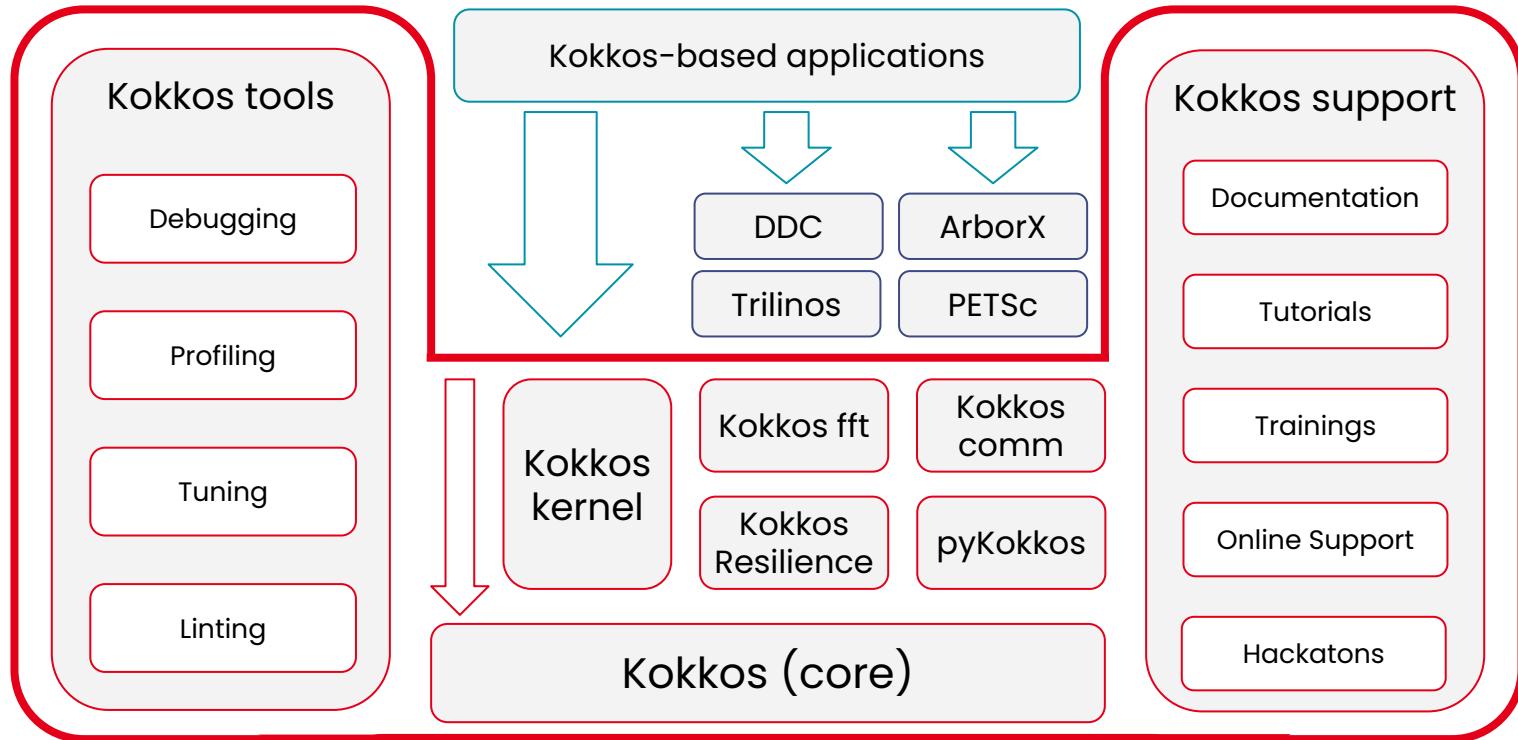


# Kokkos Ecosystem





# Kokkos Ecosystem





# To conclude



- Kokkos is a strong **vendor-neutral**, **performance portable** Exascale programming model with **GPU support**



- A **worldwide collaborative** effort for an **application first** strategy



- CExA & HPSF ensure it is a **sovereign** and **sustainable** approach that can be relied on for the foreseeable future

- A strong **dynamic** all over the CEA **and beyond**